

# Master in Artificial Intelligence

## Advanced Human Language Technologies Statistical Models of Language

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models



UNIVERSITAT POLITÈCNICA DE CATALUNYA  
BARCELONATECH

Facultat d'Informàtica de Barcelona



# Outline

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

# Outline

## 1 Statistical Models for NLP

### ■ Why modeling

#### ■ Prediction & Similarity Models

## 2 Maximum Likelihood Estimation (MLE)

### ■ Working example

### ■ Smoothing & Estimator Combination

## 3 Maximum Entropy Modeling

### ■ Overview

### ■ Building ME Models

## 4 Log-Linear Models

Statistical  
Models for  
NLP

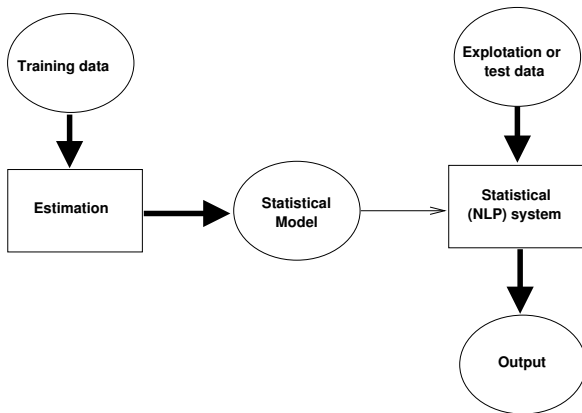
Why modeling

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# We model to make predictions



Statistical  
Models for  
NLP

Why modeling

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Outline

## 1 Statistical Models for NLP

- Why modeling
- Prediction & Similarity Models

## 2 Maximum Likelihood Estimation (MLE)

- Working example
- Smoothing & Estimator Combination

## 3 Maximum Entropy Modeling

- Overview
- Building ME Models

## 4 Log-Linear Models

Statistical  
Models for  
NLP

Prediction &  
Similarity Models

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Prediction Models & Similarity Models

Statistical  
Models for  
NLP

Prediction &  
Similarity Models

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

- **Prediction Models**: Oriented to *predict probabilities* of *future* events, knowing *past* and *present*.
- **Similarity Models**: Oriented to compute *similarities* *between objects* (may be used to predict, *EBL*).

# Similarity Models

- Objects represented as feature-vectors, feature-sets, distribution-vectors, ...
- Used to group objects (clustering, data analysis, pattern discovery, ...)
- If classified objects are available, similarity may be used as a prediction (example-based ML techniques).
- Example: Document representation
  - Documents are represented as vectors in a high dimensional  $\mathbb{R}^n$  space.
  - Dimensions are word forms, lemmas, NEs, n-grams, ...
  - Values may be either binary or real-valued (count, frequency, ...)
  - Vector-space algebra and metrics can be used

# Prediction Models

- **Estimation**: Using data to infer information about **distributions**
  - **Parametric** / **non-parametric** estimation
  - Finding good estimators: **MLE**, **MEE**, ...
  - **Explicit** / **implicit** models
- **Classification**: **Predictions** based on past behaviour
  - **Predict** most likely **target** given **classification features** (implies independence assumptions!)
  - Granularity of equivalence classes (**bins**): **discrimination power** vs. **statistical reliability**
- In general, ML models estimate (i.e. **learn**) **conditional probability distributions**  $P(\text{target}|\text{features})$
- Many NLP tasks require a posterior search step to find the best combination of predictions.



# Prediction Models

**Example:** Noisy Channel Model (Shannon 48)



## NLP Applications

Appl.	Input	Output	$p(i)$	$p(o   i)$
MT	L word sequence	M word sequence	$p(L)$	Translation model
OCR	Actual text	Text with mistakes	prob. of language text	model of OCR errors
PoS tagging	PoS tags sequence	word sequence	prob. of PoS sequence	$p(w   t)$
Speech recog.	word sequence	speech signal	prob. of word sequence	acoustic model

Given  $\mathbf{o}$ , we want to find the most likely  $\mathbf{i}$

$$\underset{i}{\operatorname{argmax}} P(i | o) = \underset{i}{\operatorname{argmax}} P(o, i) = \underset{i}{\operatorname{argmax}} P(i)P(o | i)$$

# Finding good estimators: MLE

## Maximum Likelihood Estimation (MLE)

- Choose the alternative that maximizes the probability of the observed outcome.
- $\bar{\mu}_n$  is a MLE for  $E(X)$
- $s_n^2$  is a MLE for  $\sigma^2$
- Zipf's Laws. Data sparseness. Smoothing techniques.

$P(a, b)$	dans	en	à	sur	au-cours-de	pendant	selon	
in	0.04	0.10	0.15	0	0.08	0.03	0	0.40
on	0.06	0.25	0.10	0.15	0	0	0.04	0.60
total	0.10	0.35	0.25	0.15	0.08	0.03	0.04	1.0

# Finding good estimators: MEE

## Maximum Entropy Estimation (MEE)

- Choose the alternative that **maximizes** the **entropy** of the obtained **distribution**, maintaining the observed **probabilities**.

### Observations:

$$p(\text{en} \vee \grave{\text{a}}) = 0.6;$$

$$p((\text{en} \vee \grave{\text{a}}) \wedge \text{in}) = 0.4;$$

$$p(\text{in}) = 0.5$$

$P(a, b)$	dans	en	à	sur	au-cours-de	pendant	selon	
in	0.02	0.20	0.20	0.02	0.02	0.02	0.02	0.5
on	0.06	0.10	0.10	0.06	0.06	0.06	0.06	
total		0.6						1.0

# Outline

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Outline

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Working example

Maximum  
Entropy  
Modeling

Log-Linear  
Models

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

## Working Example: N-gram models

- Predict the next element in a sequence (e.g. next character, next word, next PoS, next stock value, ... ), given the history of previous elements:

$$P(w_n | w_1 \dots w_{n-1})$$

- Markov assumption: Only local context (of size  $n - 1$ ) is taken into account.  $P(w_i | w_{i-n+1} \dots w_{i-1})$
- bigrams, trigrams, four-grams ( $n = 2, 3, 4$ ).

*Sue swallowed the large green <?>*

- Parameter estimation (number of equivalence classes)
- Parameter reduction: stemming, semantic classes, PoS, ...

Model	Parameters
bigram	$20,000^2 = 4 \times 10^8$
trigram	$20,000^3 = 8 \times 10^{12}$
four-gram	$20,000^4 = 1.6 \times 10^{17}$

Language model sizes for a 20,000 words vocabulary

# N-gram model estimation

Estimate the probability of the target feature based on observed data. The prediction task can be reduced to having good estimations of the n-gram distribution:

$$P(w_n | w_1 \dots w_{n-1}) = \frac{P(w_1 \dots w_n)}{P(w_1 \dots w_{n-1})}$$

## ■ MLE (Maximum Likelihood Estimation)

$$P_{MLE}(w_1 \dots w_n) = \frac{C(w_1 \dots w_n)}{N}$$

$$P_{MLE}(w_n | w_1 \dots w_{n-1}) = \frac{C(w_1 \dots w_n)}{C(w_1 \dots w_{n-1})}$$

- No probability mass for unseen events
- Data sparseness, Zipf's Law
- Unsuitable for NLP (widely used, though)

# Brief Parenthesis: Zipf's Laws

## Zipf's Laws (1929)

- Word frequency is inversely proportional to its rank (speaker/hearer minimum effort)  $f \sim 1/r$
- Number of senses is proportional to frequency root  $m \sim \sqrt{f}$
- Frequency of intervals between repetitions is inversely proportional to the length of the interval  $F \sim 1/I$
- Frequency based approaches are hard, since most words are rare
  - Most common 5% words account for about 50% of a text
  - 90% least common words account for less than 10% of the text
  - Almost half of the words in a text occur only once



# Outline

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Smoothing &  
Estimator  
Combination

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Notation

- $C(w_1 \dots w_n)$ : Observed occurrence count for n-gram  $w_1 \dots w_n$ .

- $N$ : Number of observed n-gram occurrences

$$N = \sum_{w_1 \dots w_n} C(w_1 \dots w_n)$$

- $N_k$ : Number of classes (n-grams) observed k times.
- $B$ : Number of equivalence classes or bins (number of potentially observable n-grams).

# Smoothing 1 - Adding Counts



- **Laplace's Law** (adding one)

$$P_{\text{LAP}}(w_1 \dots w_n) = \frac{C(w_1 \dots w_n) + 1}{N + B}$$

- For large values of B too much probability mass is assigned to unseen events

- **Lidstone's Law**

$$P_{\text{LID}}(w_1 \dots w_n) = \frac{C(w_1 \dots w_n) + \lambda}{N + B\lambda}$$

- Usually  $\lambda = 0.5$ , *Expected Likelihood Estimation*.
- Equivalent to linear interpolation between MLE and uniform prior, with  $\mu = N/(N + B\lambda)$ ,

$$P_{\text{LID}}(w_1 \dots w_n) = \mu \frac{C(w_1 \dots w_n)}{N} + (1 - \mu) \frac{1}{B}$$

## Smoothing 2 - Discounting Counts

### ■ Absolute Discounting

$$P_{\text{ABS}}(w_1 \dots w_n) = \begin{cases} \frac{C(w_1 \dots w_n) - \delta}{N} & \text{if } C(w_1 \dots w_n) > 0 \\ \frac{(B - N_0)\delta / N_0}{N} & \text{otherwise} \end{cases}$$

### ■ Linear Discounting

$$P_{\text{LIN}}(w_1 \dots w_n) = \begin{cases} (1 - \alpha) \frac{C(w_1 \dots w_n)}{N} & \text{if } C(w_1 \dots w_n) > 0 \\ \alpha / N_0 & \text{otherwise} \end{cases}$$

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Smoothing &  
Estimator  
Combination

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Combining Estimators

## ■ Simple Linear Interpolation

$$\begin{aligned} P_{LI}(w_n | w_{n-2}, w_{n-1}) &= \lambda_1 P_1(w_n) \\ &\quad + \lambda_2 P_2(w_n | w_{n-1}) \\ &\quad + \lambda_3 P_3(w_n | w_{n-2}, w_{n-1}) \end{aligned}$$

## ■ Backing-off

$$P_{BO}(w_i | h) = \begin{cases} (1 - \alpha_h) \frac{C(h, w_i)}{C(h)} & \text{if } C(h, w_i) > k \\ \delta_{h'} P_{BO}(w_i | h') & \text{otherwise} \end{cases}$$

(where  $h = w_{i-n+1} \dots w_{i-1}$ ,  $h' = w_{i-n+2} \dots w_{i-1}$ )

Different options to determine  $\alpha_h$  and  $\delta_{h'}$  (e.g.  $\alpha_h = \delta_{h'}$ ,  $\forall h$ )

# Outline

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Outline

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Overview

Log-Linear  
Models

# MEM Overview

- Maximum Entropy: alternative estimation technique.
- Able to deal with different kinds of evidence
- ME principle:
  - Do not assume anything about non-observed events.
  - Find the most uniform (maximum entropy, less informed) probability distribution that matches the observations.
- Example:

$p(x, y)$	0	1
a	?	?
b	?	?
total	0.6	1.0

*Observations*

$p(x, y)$	0	1
a	0.5	0.1
b	0.1	0.3
total	0.6	1.0

*One possible  $p(x, y)$*

$p(x, y)$	0	1
a	0.3	0.2
b	0.3	0.2
total	0.6	1.0

*Max. Entropy  $p(x, y)$*



# ME Modeling

- Observed facts are constraints for the desired model  $p$ .
- Constraints take the form of feature functions:

$$f_i : \varepsilon \rightarrow \{0, 1\}$$

- The desired model  $p$  must satisfy the constraints:  
*The expectation predicted by model  $p$  for any feature  $f_i$  must match the observed expectation for  $f_i$*   
i.e.:

$$\begin{aligned} E_p(f_i) &= E_{\tilde{p}}(f_i) \quad \forall i \\ \sum_{x \in \varepsilon} p(x) f_i(x) &= \sum_{x \in \varepsilon} \tilde{p}(x) f_i(x) \quad \forall i \end{aligned}$$

# Example

- Example:

$\varepsilon = \{a, b\} \times \{0, 1\}$	$p(x, y)$	0	1
	a	?	?
	b	?	?
	total	0.6	1.0

- Observed fact:  $p(a, 0) + p(b, 0) = 0.6$
- Encoded as a constraint:  $E_p(f_1) = 0.6$   
where:

- $f_1(x, y) = \begin{cases} 1 & \text{if } y = 0 \\ 0 & \text{otherwise} \end{cases}$

- $E_p(f_1) = \sum_{(x,y) \in \{a,b\} \times \{0,1\}} p(x, y) f_1(x, y)$

# Outline

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling  
Building ME Models

Log-Linear  
Models

# Probability Model

- There is an infinite set  $P$  of probability models consistent with observations:

$$P = \{p \mid E_p(f_i) = E_{\tilde{p}}(f_i), \forall i\}$$

- Maximum entropy model

$$\begin{aligned} p^* &= \operatorname{argmax}_{p \in P} H(p) \\ &= \operatorname{argmax}_{p \in P} \left( - \sum_{x \in \mathcal{E}} p(x) \log p(x) \right) \end{aligned}$$

# Conditional Probability Model

- For NLP applications, we are usually interested in conditional distributions  $P(Y|X)$ , thus, the ME model is

$$p^* = \operatorname{argmax}_{p \in P} H(p) = \operatorname{argmax}_{p \in P} H(Y | X)$$

where:

$$\begin{aligned} H(Y | X) &= \sum_{x \in X} p(x) H(Y | X = x) \\ &= - \sum_{x \in X} p(x) \sum_{y \in Y} p(y | x) \log p(y | x) \\ &= - \sum_{x \in X, y \in Y} p(x, y) \log p(y | x) \\ &= - \sum_{x \in X, y \in Y} p(x, y) \log \frac{p(x, y)}{p(x)} \end{aligned}$$

# Parameter Estimation

Example: Maximum entropy model for translating *in* to French

- No constraints

$P(x)$	dans	en	à	au-cours-de	pendant	
	0.2	0.2	0.2	0.2	0.2	
total						1.0

- With constraint  $p(\text{dans}) + p(\text{en}) = 0.3$

$P(x)$	dans	en	à	au-cours-de	pendant	
	0.15	0.15	0.233	0.233	0.233	
total	<b>0.3</b>					1.0

- With constraints

$$p(\text{dans}) + p(\text{en}) = 0.3; \quad p(\text{en}) + p(\text{à}) = 0.5$$

...Not so easy !

# Parameter estimation

- Exponential models

$$p(y | x) = \frac{1}{Z(x)} \prod_{j=1}^k \alpha_j^{f_j(x,y)} \quad \alpha_j > 0, \quad Z(x) = \sum_y \prod_{i=1}^k \alpha_i^{f_i(x,y)}$$

- Can also be formulated as

$$p(y | x) = \frac{1}{Z(x)} \exp \left( \sum_{j=1}^k \lambda_j f_j(x, y) \right) \quad (\text{i.e. } \lambda_i = \ln \alpha_i)$$

- Each model parameter weights the influence of a feature.
- Optimal parameters can be computed with:
  - Generalized Iterative Scaling (GIS) [Darroch & Ratcliff 72]
  - Improved Iterative Scaling (IIS) [Della Pietra et al. 96]
  - Limited Memory BFGS (**LM-BFGS**) [Malouf 03]

# Example: Text Categorization

- Probabilistic model over  $W \times C$  (Words  $\times$  Categories).

A document is a set of words:  $d = (w_1, w_2 \dots w_N)$ .

Each combination  $w, c \in W \times C$  is a feature:

$$f_{w,c}(d, c') = \begin{cases} \frac{N(w,d)}{N(d)} & \text{if } c = c' \\ 0 & \text{otherwise} \end{cases}$$

- Disambiguation: Select class with highest  $P(c \mid d)$

$$P(c \mid d) = \frac{1}{Z(d)} \exp\left(\sum_i \lambda_i f_i(d, c)\right)$$



# MEM Summary

- Advantages

- Teoretically well founded
- Enables combination of random context features
- Better probabilistic models than MLE (no smoothing needed)
- General approach (features, events and classes)

- Disadvantages

- Implicit probabilistic model (joint or conditional probability distribution obtained from model parameters).

ME Models are a particular case of **Log-Linear models**

# Outline

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Maximum Entropy Modeling
  - Overview
  - Building ME Models
- 4 Log-Linear Models

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Log-Linear Models

$$P(y \mid x; \mathbf{w}) = \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{\sum_y \exp(\mathbf{w} \cdot \mathbf{f}(x, y))}$$

where

- $\mathbf{f}(x, y)$  is a feature vector representing  $x$  and  $y$
- $\mathbf{w}$  are the parameters of the model
- $\mathbf{w} \cdot \mathbf{f}(x, y)$  is a score for  $x$  and  $y$
- $Z(x) = \sum_y \exp(\mathbf{w} \cdot \mathbf{f}(x, y))$  is a normalizer (sums over all possible values  $y$  for  $x$ ); it's sometimes called the *partition function*

# Features, Indicator Features

- $\mathbf{f}(x, y)$  is a vector of  $d$  features representing  $x$  and  $y$

$$(f_1(x, y), \dots, f_j(x, y), \dots, f_d(x, y))$$

- What's in a feature  $f_j(x, y)$ ?

- Anything we can compute using  $x$  and  $y$
- Anything that is **informative** for (or **against**)  $x$  belonging to class  $y$
- **Indicator features**: binary-valued features looking at a single simple property

$$f_j(c, b) = \begin{cases} 1 & \text{if prefix}(c) = Mr \text{ and } b = \text{no} \\ 0 & \text{otherwise} \end{cases}$$

$$f_k(c, b) = \begin{cases} 1 & \text{if uppercase(next}(c)) \text{ and } b = \text{yes} \\ 0 & \text{otherwise} \end{cases}$$

# Features, Parameters, Inner Products

$$P(y \mid x; \mathbf{w}) = \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{\sum_y \exp(\mathbf{w} \cdot \mathbf{f}(x, y))}$$

- $\mathbf{f}(x, y) \in \mathbb{R}^d$  is a feature vector with  $d$  features
- $\mathbf{w} \in \mathbb{R}^d$  is a **parameter** vector, with  $d$  parameters
- Inner products (*a.k.a.* **dot** products)

$$\mathbf{w} \cdot \mathbf{f}(x, y) = \sum_{i=1}^d \mathbf{w}_i \mathbf{f}_i(x, y)$$

# Log-linear Models

$$P(y \mid x; \mathbf{w}) = \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{\sum_y \exp(\mathbf{w} \cdot \mathbf{f}(x, y))}$$

where

- $\mathbf{f}(x, y)$  is a feature vector representing  $x$  and  $y$ 
  - **Arbitrary** features of  $x$  and  $y$  are allowed
  - They are **provided** for the application in turn
- $\mathbf{w}$  are the parameters of the model
- **Two problems:**
  - How to make **predictions** using  $P(y \mid x)$
  - How to **estimate** the parameters  $\mathbf{w}$ ?

# Log-linear Models: Name

- Let's take the **log** of the conditional probability:

$$\begin{aligned}\log P(y \mid x; \mathbf{w}) &= \log \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{\sum_y \exp(\mathbf{w} \cdot \mathbf{f}(x, y))} \\ &= \mathbf{w} \cdot \mathbf{f}(x, y) - \log \sum_y \exp(\mathbf{w} \cdot \mathbf{f}(x, y)) \\ &= \boxed{\mathbf{w} \cdot \mathbf{f}(x, y)} - \log Z(x)\end{aligned}$$

- Partition function:  $Z(x) = \sum_y \exp(\mathbf{w} \cdot \mathbf{f}(x, y))$
- $\log Z(x)$  is a **constant** for a fixed  $x$
- In the **log** space, computations are **linear**

# Log-linear Models: Making Predictions

- Given  $x$ , what  $y$  in  $\{1, \dots, L\}$  is most appropriate?

$$\begin{aligned}\text{best}(x) &= \operatorname{argmax}_{y \in \{1, \dots, L\}} P(y \mid x; \mathbf{w}) \\ &= \operatorname{argmax}_{y \in \{1, \dots, L\}} \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{Z(x)} \\ &= \operatorname{argmax}_{y \in \{1, \dots, L\}} \exp(\mathbf{w} \cdot \mathbf{f}(x, y)) \\ &= \operatorname{argmax}_{y \in \{1, \dots, L\}} \mathbf{w} \cdot \mathbf{f}(x, y)\end{aligned}$$

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models



# Log-linear Models: Making Predictions

- Given  $x$ , what  $y$  in  $\{1, \dots, L\}$  is most appropriate?

$$\begin{aligned}\text{best}(x) &= \underset{y \in \{1, \dots, L\}}{\operatorname{argmax}} P(y \mid x; \mathbf{w}) \\ &= \underset{y \in \{1, \dots, L\}}{\operatorname{argmax}} \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{Z(x)} \\ &= \underset{y \in \{1, \dots, L\}}{\operatorname{argmax}} \exp(\mathbf{w} \cdot \mathbf{f}(x, y)) \\ &= \underset{y \in \{1, \dots, L\}}{\operatorname{argmax}} \mathbf{w} \cdot \mathbf{f}(x, y)\end{aligned}$$

- Predictions only require simple inner products (linear)
- No need to exponentiate!

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Log-linear Models: Computing Probabilities

$$P(y \mid x; \mathbf{w}) = \frac{\exp(\mathbf{w} \cdot \mathbf{f}(x, y))}{Z(x)}$$

- Sometimes we will be interested in computing  $P(y \mid x)$

- It can be used as a **measure of confidence**, e.g.

$$P(\text{yes} \mid c) = 0.51 \text{ versus}$$

$$P(\text{yes} \mid c) = 0.99$$

- We need to compute:

$$Z(x) = \sum_{y=\{1,\dots,L\}} \exp(\mathbf{w} \cdot \mathbf{f}(x, y))$$

- Fast as long as  $L$  is **not too large**

# Parameter Estimation in Log-linear Models

- How to **estimate** model **parameters**  $\mathbf{w}$  given a training set:

$$\left\{ (x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)}) \right\}$$

- Let's define the conditional log-likelihood of the data:

$$L(\mathbf{w}) = \frac{1}{m} \sum_{k=1}^m \log P(y^{(k)} | x^{(k)}; \mathbf{w})$$

- $L(\mathbf{w})$  measures how well  $\mathbf{w}$  explains the data. A good value for  $\mathbf{w}$  will give a high value for  $P(y^{(k)} | x^{(k)}; \mathbf{w})$  for all  $k = 1 \dots m$ .
- We want  $\mathbf{w}$  that **maximizes**  $L(\mathbf{w})$

# Parameter Estimation in Log-Linear Models

- We pose it as an **optimization** problem
- Find:

$$\mathbf{w}^* = \underset{\mathbf{w} \in \mathbb{R}^d}{\operatorname{argmax}} L(\mathbf{w})$$

where

- But low-frequency features may end up having large weights (i.e. **overfitting**)
- We need a **regularization** factor that penalizes solutions with a large norm (similar to norm-minimization in SVM):

$$L'(\mathbf{w}) = \frac{1}{m} \sum_{k=1}^m \log P(y^{(k)} | x^{(k)}; \mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- where  $\lambda$  is a parameter to control the trade-off between fitting the data and model complexity. **Tuned experimentally.**

# Parameter Estimation in Log-Linear Models

- So we want to find:

$$\begin{aligned}\mathbf{w}^* &= \operatorname{argmax}_{\mathbf{w} \in \mathbb{R}^d} L'(\mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{w} \in \mathbb{R}^d} \left( \frac{1}{m} \sum_{k=1}^m \log P(\mathbf{y}^{(k)} | \mathbf{x}^{(k)}; \mathbf{w}) - \frac{\lambda}{2} \|\mathbf{w}\|^2 \right)\end{aligned}$$

- In general there is no analytical solution to this optimization
- ... but it is a **convex** function  $\Rightarrow$  We use iterative techniques, i.e. gradient-based optimization
- Very fast algorithms exist (e.g. **LBFGS**)

# Parameter Estimation in Log-Linear Models : Gradient step

- Initialize  $\mathbf{w} = \mathbf{0}$
- Repeat
  - Compute gradient  $\delta = (\delta_1, \dots, \delta_d)$ , where:

$$\delta_j = \frac{\partial L'(\mathbf{w})}{\partial \mathbf{w}_j} \quad \forall j = 1 \dots d$$

- Compute **step** size

$$\beta^* = \operatorname{argmax}_{\beta \in \mathbb{R}} L'(\mathbf{w} + \beta \delta)$$

- Move  $\mathbf{w}$  in the direction of the gradient

$$\mathbf{w} \leftarrow \mathbf{w} + \beta^* \delta$$

- until convergence ( $\|\delta\| < \epsilon$ )

# Log-linear Models: Computing the Gradient

$$\begin{aligned} \frac{\partial L'(\mathbf{w})}{\partial \mathbf{w}_j} &= \frac{1}{m} \sum_{k=1}^m \mathbf{f}_j(\mathbf{x}^{(k)}, y^{(k)}) \\ &\quad - \sum_{k=1}^m \sum_{y \in \{1, \dots, L\}} P(y|\mathbf{x}^{(k)}; \mathbf{w}) \mathbf{f}_j(\mathbf{x}^{(k)}, y) \\ &\quad - \lambda \mathbf{w}_j \end{aligned}$$

- First term: observed **mean** feature value
- Second term: **expected** feature **value** under current  $\mathbf{w}$
- In the optimal, observed = expected

**Maximum log-likelihood log-linear models  
correspond to Maximum Entropy models**

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# MEM Overview

ME Models are the **dual formulation** of **Log-Linear models**.

- Maximum Entropy: alternative estimation technique.
- Able to deal with different kinds of evidence
- ME principle:
  - Do not assume anything about non-observed events.
  - Find the most uniform (maximum entropy, less informed) probability distribution that matches the observations.
- Example:

$p(x, y)$	0	1	
a	?	?	
b	?	?	
total	0.6	1.0	

*Observations*

$p(x, y)$	0	1	
a	0.5	0.1	
b	0.1	0.3	
total	0.6	1.0	

*One possible  $p(x, y)$*

$p(x, y)$	0	1	
a	0.3	0.2	
b	0.3	0.2	
total	0.6	1.0	

*Max.Entropy  $p(x, y)$*



# ME Modeling

- Observed facts are constraints for the desired model  $p$ .
- Constraints take the form of feature functions:

$$f_i : \mathcal{E} \rightarrow \{0, 1\}$$

- The desired model  $p$  must satisfy the constraints:  
*The expectation predicted by model  $p$  for any feature  $f_i$  must match the observed expectation for  $f_i$*   
i.e.:

$$\begin{aligned} E_p(f_i) &= E_{\tilde{p}}(f_i) \quad \forall i \\ \sum_{x \in \mathcal{E}} p(x) f_i(x) &= \sum_{x \in \mathcal{E}} \tilde{p}(x) f_i(x) \quad \forall i \end{aligned}$$

# Probability Model

- There is an infinite set  $P$  of probability models consistent with observations:

$$P = \{p \mid E_p(f_i) = E_{\tilde{p}}(f_i), \forall i\}$$

- Maximum entropy model

$$\begin{aligned} p^* &= \operatorname{argmax}_{p \in P} H(p) \\ &= \operatorname{argmax}_{p \in P} \left( - \sum_{x \in \mathcal{E}} p(x) \log p(x) \right) \end{aligned}$$

# Conditional Probability Model

- For NLP applications, we are usually interested in **conditional distributions**  $P(Y|X)$ , thus, the ME model is

$$p^* = \operatorname{argmax}_{p \in \mathcal{P}} H(p) = \operatorname{argmax}_{p \in \mathcal{P}} H(Y | X)$$

where:

$$\begin{aligned} H(Y | X) &= \sum_{x \in X} p(x) H(Y | X = x) \\ &= - \sum_{x \in X} p(x) \sum_{y \in Y} p(y | x) \log p(y | x) \\ &= - \sum_{x \in X, y \in Y} p(x, y) \log p(y | x) \end{aligned}$$

# Parameter Estimation

Example: Maximum entropy model for translating *in* to French

- No constraints

$P(x)$	dans	en	à	au-cours-de	pendant	
	0.2	0.2	0.2	0.2	0.2	
total						1.0

- With constraint  $p(\text{dans}) + p(\text{en}) = 0.3$

$P(x)$	dans	en	à	au-cours-de	pendant	
	0.15	0.15	0.233	0.233	0.233	
total	<b>0.3</b>					1.0

- With constraints

$$p(\text{dans}) + p(\text{en}) = 0.3; \quad p(\text{en}) + p(\text{à}) = 0.5$$

...Not so easy !

# Parameter estimation

- ME models are exponential models, same as log-linear models

$$p(y | x) = \frac{1}{Z(x)} \exp \left( \sum_{j=1}^k \lambda_j f_j(x, y) \right)$$

$$\text{where } Z(x) = \sum_{y'} \exp \left( \sum_{j=1}^k \lambda_j f_j(x, y') \right)$$

- Each model parameter weights the influence of a feature.
- Same convex optimization algorithms are used (e.g. **LM-BFGS** [Malouf 03])

# Outline

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Log-Linear  
Models

Examples

- 1 Statistical Models for NLP
  - Why modeling
  - Prediction & Similarity Models
- 2 Maximum Likelihood Estimation (MLE)
  - Working example
  - Smoothing & Estimator Combination
- 3 Log-Linear Models
  - Maximum Entropy Models
  - Examples

# Example: Text Categorization

- Probabilistic model over  $W \times C$  (Words  $\times$  Categories).

A document is a set of words:  $d = (w_1, w_2 \dots w_N)$ .

Each combination  $w, c \in W \times C$  is a feature:

$$f_{w,c}(d, c') = \begin{cases} \frac{N(w,d)}{N(d)} & \text{if } c = c' \\ 0 & \text{otherwise} \end{cases}$$

- Disambiguation: Select class with highest  $P(c \mid d)$

$$P(c \mid d) = \frac{1}{Z(d)} \exp\left(\sum_i \lambda_i f_i(d, c)\right)$$

## Example: Identifying Sentence Boundaries

*The president lives in Washington,  
D.C. The presidents met in Washing-  
ton D.C. in 2010. Mr. Wayne is young.  
Mr. Wayne is a Ph.D. I got 98.5%!  
What?*

**Goal:** given a text, identify tokens that end a sentence

- Candidate characters: . ! ?
- Candidate tokens: tokens containing candidate characters
- Given a candidate token in a *context* decide whether it ends a sentence or not



# Example: Sentence Boundaries

- Candidate: punctuation sign + context  
 $c = \langle \text{sign, prefix, suffix, previous, next} \rangle$
- Assume access to *annotated* data:

b	sign	prefix	suffix	prev	next
-----					
no	.	D	C.	Washington,	The
yes	.	D.C		Washington,	The
no	.	Mr		2010.	Wayne

- Let's take a probabilistic approach:
  - $P(\text{yes} | c)$ : conditional probability of  $c$  being end of sentence
  - $P(\text{no} | c)$ : conditional probability of  $c$  *not* being e.o.s.
  - Obviously,  $P(\text{yes} | c) + P(\text{no} | c) = 1$
  - Predict yes if  $P(\text{yes} | c) > 0.5$
- How to model  $P(\text{yes} | c)$  and  $P(\text{no} | c)$ ?

# Example system: Identifying Sentence Boundaries

(Reynar and Ratnaparkhi '97)

- Candidate: punctuation sign + context  
 $c = \langle \text{sign, prefix, suffix, previous, next} \rangle$
- **Goal:** estimate  $P(\text{yes} \mid c)$  and  $P(\text{no} \mid c)$
- Feature templates:
  - 1 The prefix
  - 2 The suffix
  - 3 The word previous
  - 4 The word next
  - 5 Whether prefix or suffix are in ABBREVIATIONS
    - ABBREVIATIONS: list of all training tokens that contain a . and are *not* sentence boundaries
  - 6 Whether previous or next are in ABBREVIATIONS
- Actual features are generated by applying each template to each training example

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

# Example system: Identifying Sentence Boundaries

(Reynar and Ratnaparkhi '97)

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

## FEATURE TEMPLATES

- 1 The prefix
- 2 The suffix
- 3 The word previous
- 4 The word next
- 5 Whether prefix or suffix are in ABBREVIATIONS
- 6 Whether previous or next are in ABBREVIATIONS

< b=no punc=. pref=Mr suff= prev=2010. next=Wayne >

## GENERATED FEATURES

$$\begin{aligned} f_1(c, b) &= \begin{cases} 1 & \text{if pref}(c) = \text{Mr} \\ & \text{and } b = \text{no} \\ 0 & \text{otherwise} \end{cases} & f_4(c, b) &= \begin{cases} 1 & \text{if next}(c) = \text{Wayne} \\ & \text{and } b = \text{no} \\ 0 & \text{otherwise} \end{cases} \\ f_2(c, b) &= \begin{cases} 1 & \text{if suff}(c) = \text{NULL} \\ & \text{and } b = \text{no} \\ 0 & \text{otherwise} \end{cases} & f_5(c, b) &= \begin{cases} 1 & \text{if (abbr(pref}(c)) \text{ or abbr(suff}(c)))} \\ & \text{and } b = \text{no} \\ 0 & \text{otherwise} \end{cases} \\ f_3(c, b) &= \begin{cases} 1 & \text{if prev}(c) = \text{2010.} \\ & \text{and } b = \text{no} \\ 0 & \text{otherwise} \end{cases} & f_6(c, b) &= \begin{cases} 1 & \text{if (abbr(prev}(c)) \text{ or abbr(next}(c)))} \\ & \text{and } b = \text{no} \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

# Example System: Identifying Sentence Boundaries

(Reynar and Ratnaparkhi '97)

Statistical  
Models for  
NLP

Maximum  
Likelihood  
Estimation  
(MLE)

Maximum  
Entropy  
Modeling

Log-Linear  
Models

training sentences	test accuracy
500	96.5%
1000	97.3%
2000	97.3%
4000	97.6%
8000	97.6%
16000	97.8%
39441	98.0%

- Corpus: Wall Street Journal, English